

🛡️ AI-Powered Phishing Email Detector

An end-to-end machine learning system that detects phishing emails using natural language processing and classification algorithms.

Python 3.8+

License MIT

ML Scikit-Learn

⌚ Table of Contents

- [Overview](#)
- [Features](#)
- [Project Structure](#)
- [Dataset](#)
- [Installation](#)
- [Usage](#)
- [Model Performance](#)
- [Architecture](#)
- [Ethical Considerations](#)
- [Contributing](#)
- [License](#)

🎯 Overview

Phishing emails remain one of the most common initial attack vectors for cybercriminals. Traditional rule-based systems often fail against new, adaptive phishing techniques. This project implements a machine learning approach to detect phishing emails by analyzing text patterns, linguistic features, and structural characteristics.

PROF

Problem Statement

- Phishing attacks cause billions of dollars in losses annually
- Rule-based detection systems can't adapt to new phishing techniques
- Need for automated, intelligent email classification

Solution

A Logistic Regression classifier trained on TF-IDF features from email text, providing:

- Binary classification (phishing vs. legitimate)
- Probability scores for risk assessment
- Fast, explainable predictions

✨ Features

- **Data Pipeline:** Automated data ingestion and cleaning
- **Text Preprocessing:** URL/email normalization, HTML removal, tokenization
- **Feature Engineering:** TF-IDF vectorization with n-grams
- **Model Training:** Logistic Regression with cross-validation
- **Evaluation:** Comprehensive metrics and visualizations
- **Web Interface:** Interactive Streamlit application for real-time predictions

📁 Project Structure

```

ai-phishing-email-detector/
├── data/
│   ├── raw/                                # Raw email datasets
│   └── processed/                           # Cleaned and processed data
├── docs/
│   ├── evaluation_report.txt               # Model evaluation results
│   └── plots/                               # Visualization outputs
├── notebooks/                            # Jupyter notebooks for exploration
└── src/
    ├── data/
    │   └── make_dataset.py                  # Data ingestion module
    ├── features/
    │   └── build_features.py              # Feature engineering module
    ├── models/
    │   ├── train.py                      # Model training module
    │   └── evaluate.py                  # Model evaluation module
    ├── ui/
    │   └── app.py                       # Streamlit web application
    ├── api/
    │   └── utils/                        # REST API (optional)
    ├── tests/                            # Utility functions
    ├── .gitignore
    ├── README.md
    └── requirements.txt

```

PROF

📊 Dataset

Recommended Datasets

Dataset	Source	Description
Enron Spam	Kaggle	Classic email spam dataset
Phishing Email	Kaggle	Labeled phishing emails
CEAS 2008	CEAS	Conference anti-spam corpus

Data Format

The dataset should be a CSV file with at least these columns:

```
email_text, label  
"Email content here...", 0  
"Suspicious email here...", 1
```

Where:

- `label = 0` → Legitimate email
- `label = 1` → Phishing email

Data Placement

Place your dataset at: `data/raw/emails.csv`

 **Note:** Raw data is excluded from version control. Check dataset licenses before redistribution.

Installation

Prerequisites

- Python 3.8 or higher
- pip package manager
- Git

Setup

1. Clone the repository

```
git clone https://github.com/<your-username>/ai-phishing-email-detector.git  
cd ai-phishing-email-detector
```

PROF

2. Create virtual environment

```
python -m venv .venv  
source .venv/bin/activate # Linux/Mac  
# or  
.venv\Scripts\activate # Windows
```

3. Install dependencies

```
pip install -U pip  
pip install -r requirements.txt
```

4. Download dataset

- Download a phishing email dataset from Kaggle
- Place it at `data/raw/emails.csv`

💻 Usage

1. Data Ingestion

Clean and preprocess the raw email data:

```
python src/data/make_dataset.py
```

2. Feature Engineering

Extract TF-IDF features from email text:

```
python src/features/build_features.py
```

3. Model Training

Train the Logistic Regression classifier:

```
python src/models/train.py
```

4. Model Evaluation

Generate evaluation metrics and plots:

—
PROF

```
python src/models/evaluate.py
```

5. Run Web Application

Launch the interactive Streamlit interface:

```
streamlit run src/ui/app.py
```

Then open <http://localhost:8501> in your browser.

↵ Model Performance

Classification Report

	precision	recall	f1-score	support
Legitimate	0.XX	0.XX	0.XX	XXXX
Phishing	0.XX	0.XX	0.XX	XXXX
accuracy			0.XX	XXXX
macro avg	0.XX	0.XX	0.XX	XXXX
weighted avg	0.XX	0.XX	0.XX	XXXX

Key Metrics

Metric	Value
Accuracy	XX.XX%
Precision	XX.XX%
Recall	XX.XX%
F1-Score	XX.XX%
ROC-AUC	XX.XX%

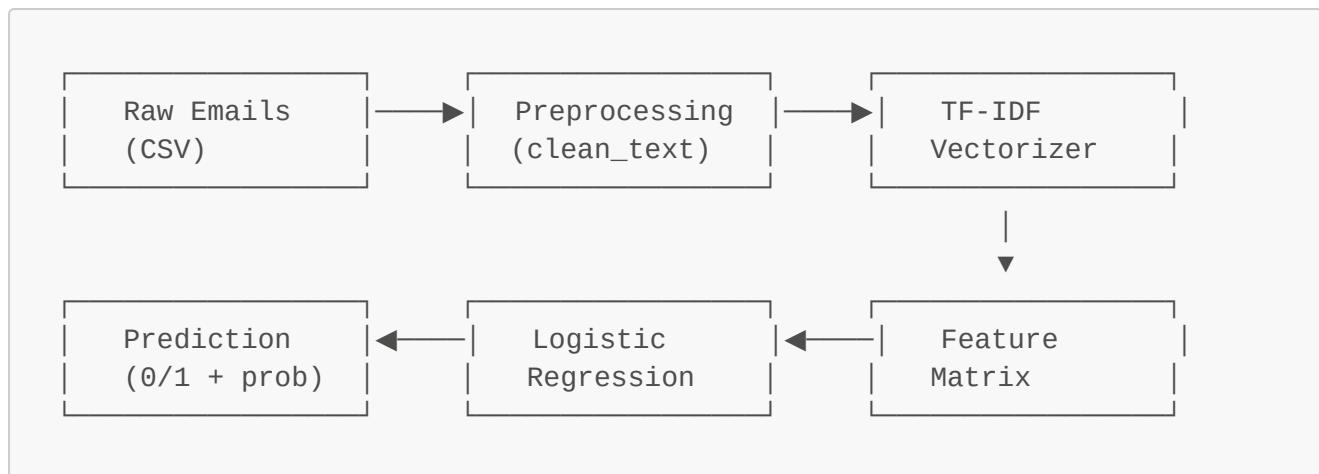
Confusion Matrix

		Predicted	
		Leg	Phi
Actual	Leg	TN	FP
	Phi	FN	TP

PROF

Note: Replace XX values with actual metrics after training.

Architecture



Pipeline Steps

1. **Data Ingestion:** Load and validate raw email data
2. **Preprocessing:** Clean text, normalize URLs/emails, remove HTML
3. **Feature Extraction:** TF-IDF vectorization with unigrams and bigrams
4. **Model Training:** Logistic Regression with balanced class weights
5. **Evaluation:** Classification metrics, confusion matrix, ROC curve
6. **Inference:** Real-time prediction via Streamlit UI

⚖️ Ethical Considerations

Privacy

- Only public datasets are used
- No real private emails are processed
- No personal data is stored or transmitted

Bias and Fairness

- Training data may have inherent biases
- Model performance may vary across different email types
- Regular evaluation on diverse datasets is recommended

Misuse Prevention

- This tool is for **educational and research purposes**
- Should not be the sole basis for blocking emails
- False positives can cause legitimate emails to be missed

Responsible AI

- Model decisions are explainable (feature coefficients)
- Probability scores provided for risk assessment
- Users encouraged to verify through official channels

—
PROF

Limitations

- May not detect zero-day phishing techniques
- Performance depends on training data quality
- Text-only analysis (doesn't check attachments or links)

🤝 Contributing

Contributions are welcome! Please follow these steps:

1. Fork the repository
2. Create a feature branch (`git checkout -b feature/amazing-feature`)
3. Commit your changes (`git commit -m 'Add amazing feature'`)
4. Push to the branch (`git push origin feature/amazing-feature`)

5. Open a Pull Request

Development Setup

```
# Install development dependencies  
pip install -r requirements.txt  
  
# Run tests  
pytest tests/  
  
# Format code  
black src/  
  
# Lint code  
flake8 src/
```

License

This project is licensed under the MIT License - see the [LICENSE](#) file for details.

Acknowledgments

- [Scikit-learn](#) for ML algorithms
- [Streamlit](#) for the web interface
- [Kaggle](#) for public datasets
- The cybersecurity community for research and datasets

⚠ Disclaimer: This tool is for educational purposes only. Always verify suspicious emails through official channels and consult cybersecurity professionals for critical decisions.